

CV Homework 4 Report

— 106062202 陳騰鴻

— Dec 28, 2019

- CV Homework 4 Report
 - Q1: Binary Classification
 - Q1-1
 - Q1-2
 - Q1-3
 - Q1-4
 - Q1-5
 - Q1-6
 - Q1-7
 - Q1-8
 - Q1-9
 - Q2: Semantic Segmentation
 - Q2-1
 - Q2-2
 - Q2-3
 - Q2-4

Q1: Binary Classification

Q1-1

```
Training complete in 4m 15s
Best val Acc: 0.830000
```

Q1-2

```
Training complete in 5m 24s
Best val Acc: 0.875000
```

Q1-3

```
Training complete in 5m 26s
Best val Acc: 0.717500
```

Q1-4

Comparisons

Model	Pretrained	Backprop Parameters	Performance
Feature Extractor	Yes	1 Linear Classifier	Second
Fine-tuned	Yes	All Layers + 1 Linear Classifier	Best
AlexNet	No	All layers	Last

Why Feature Extractor Performs Better than AlexNet:

The feature extractor benefits from the pretrained parameters, which are trained on an extensive amount of images.

Due to the effect of transfer learning, the feature extractor yields better performance.

Why Fine-tuned Model Performs Better than Feature Extractor:

2. Model: EfficientNet-B2, loaded from [rwightman/gen-efficientnet-pytorch](https://github.com/rwightman/gen-efficientnet-pytorch)

(<https://github.com/rwightman/gen-efficientnet-pytorch>).

The model has a one-layer classifier with output feature = 1000, so `nn.Linear(1000, 2)` is added.

```
1 model_ft = torch.hub.load('rwightman/gen-efficientnet-pytorch',
2                           'efficientnet_b2', pretrained=True)
3 model_ft.classifier = nn.Sequential(model_ft.classifier,
4                                     nn.Linear(1000, 2))
```

3. Optimizer: Adam with learning rate = 0.001.

```
1 optimizer_ft = optim.Adam(model_ft.parameters(), lr=0.001)
```

Q1-9

Result

```
Training complete in 14m 12s
Best val Acc: 0.907500
```

How I Chose The Modifications

1. `transforms.ColorJitter` is a commonly used transformation.
2. Auto-augmentation is one of the most implemented (<https://paperswithcode.com/task/image-augmentation>) image augmentation techniques. Augmentation for CIFAR_10 is used in hope of increasing the performance.
3. EfficientNet is chosen because of its great performance (<https://paperswithcode.com/sota/image-classification-on-imagenet>) and relatively small number of parameters. Despite EfficientNet-B7 having similar number of parameters as AlexNet does, it takes way too long to train. The final decision was to use EfficientNet-B2, only 9.2 M parameters (0.15x of AlexNet).
4. Optimizer was changed to Adam in attempt to increase the speed of convergence.

Q2: Semantic Segmentation

Q2-1

Before Removing Skip Connections

```
The highest mIOU is 0.4295 and is achieved at epoch-41
The highest pixel accuracy is 0.8523 and is achieved at epoch-41
```

After Removing Skip Connections

```
The highest mIOU is 0.4055 and is achieved at epoch-20
The highest pixel accuracy is 0.8384 and is achieved at epoch-17
```

Q2-2

Implementation

I deleted the addition of `x4` and `x3` in `FCN8s.forward`, which are the two skip connections in the neural network.

```

1 | def forward(self, x):
2 |     output = self.pretrained_net(x)
3 |     x5 = output['x5']
4 |
5 |     score = self.relu(self.deconv1(x5))
6 |     score = self.bn1(score)
7 |     score = self.relu(self.deconv2(score))
8 |     score = self.bn2(score)
9 |     score = self.bn3(self.relu(self.deconv3(score)))
10 |    score = self.bn4(self.relu(self.deconv4(score)))
11 |    score = self.bn5(self.relu(self.deconv5(score)))
12 |    score = self.classifier(score)
13 |
14 |    return score

```

Are skip connections quantitatively beneficial?

According to the experimental results, the neural network with skip connections performs better in both measurements by 2%.

Therefore, skip connections are quantitatively beneficial.

Why adding skip connections improves performance

Deep features can be obtained when data is passed through more layers.

However, spacial location information is also lost due to convolution. (Paraphrasing a quote from [this article](https://towardsdatascience.com/review-fcn-semantic-segmentation-eb8c9b50d2d1) (https://towardsdatascience.com/review-fcn-semantic-segmentation-eb8c9b50d2d1).)

Hence, by adding skip connections, the spacial location information (Extracted when the network is still shallow) is preserved and improves the performance.

Q2-3

Results

```

The highest mIOU is 0.8261 and is achieved at epoch-20
The highest pixel accuracy is 0.9629 and is achieved at epoch-20

```

Implementation

- Number of segmentation classes is changed to 3.

```

1 | num_class = 3

```

- Labels of each class can be found in `save_result_comparison`.

```

1 | if output_np[i,j] == 0: # Sky
2 |     im_seg_RGB[i,j,:] = [128, 128, 128]
3 | elif output_np[i,j] == 1: # Building
4 |     im_seg_RGB[i,j,:] = [128, 0, 0]
5 | elif output_np[i,j] == 2: # Pole
6 |     im_seg_RGB[i,j,:] = [192, 192, 128]
7 | elif output_np[i,j] == 3: # Road
8 |     im_seg_RGB[i,j,:] = [128, 64, 128]
9 | elif output_np[i,j] == 4: # Pavement
10 |    im_seg_RGB[i,j,:] = [0, 0, 192]
11 | elif output_np[i,j] == 5: # Tree
12 |    im_seg_RGB[i,j,:] = [128, 128, 0]
13 | elif output_np[i,j] == 6: # Sign Symbol
14 |    im_seg_RGB[i,j,:] = [192, 128, 128]
15 | elif output_np[i,j] == 7: # Fence
16 |    im_seg_RGB[i,j,:] = [64, 64, 128]
17 | elif output_np[i,j] == 8: # Car
18 |    im_seg_RGB[i,j,:] = [64, 0, 128]
19 | elif output_np[i,j] == 9: # Pedestrian
20 |    im_seg_RGB[i,j,:] = [64, 64, 0]
21 | elif output_np[i,j] == 10: # Bicyclist
22 |    im_seg_RGB[i,j,:] = [0, 128, 192]

```

- Labels are changed according to the specs, using boolean indexing.

```

1 | label = np.array(label_image)
2 |
3 | # Class 1
4 | mask_class1 = (label > 0) & (label < 8) | (label > 10)
5 | label[mask_class1] = 1
6 |
7 | # Class 2
8 | mask_class2 = (label > 7) & (label < 11)
9 | label[mask_class2] = 2

```

Q2-4

The Comparison Model

Q2-3. Please try to **FURTHER** reduce the number of classes from 11 to 3 ...

Thus, our comparison model should be the one **without** skip connections.

Why The Mean IoU And The Accuracy Increased Dramatically

- With less classes, the ground truth segmentations would have simpler contours, thus making it an easier task for a model to learn.
- The following are the IoUs at maximum performance:

Classifier	Result
11-class	<code>pix_acc: 0.8306, meanIoU: 0.3932, IoUs: [0.8899 0.7452 0. 0.9159 0.7039 0.738 0. 0.0034 0.2936 0.0353 0.]</code>
3-class	<code>pix_acc: 0.9629, meanIoU: 0.8261, IoUs: [0.9049 0.9559 0.6175]</code>

As we can see, the 11-class classifier fails to detect class 10, which greatly reduces the mean IoU.

- As all the other factors (hyperparameters, data augmentation, training epochs ...) are controlled, the performance increases.